

AI WRITING WORKSPACE

Write deeper.
Remember everything.


Lagoon is a local-first AI writing workspace built for long-form collaborative fiction — character-driven, context-aware, and designed around how writers actually work.

FLASK · VANILLA JS

VENICE · OLLAMA · ANY OPENAI ENDPOINT

RUNS LOCALLY

NO SUBSCRIPTION



The screenshot displays the Lagoon AI Writing Workspace interface. On the left, a sidebar shows a list of characters: Hermione Granger, Kelly Thompson, Lagoon Assistant, Leela, and Nathan Young. The main chat area shows a conversation with Amy. The chat history includes several paragraphs of text, with the latest message from Amy being: "Okay," Amy said. "Breathe, Kel. Just breathe." She paused. "You're eighteen. That matters." The chat area also shows a list of system prompts and user-defined instructions. On the right, a panel shows the current model (GPT-4o) and a list of tools (mum, headphones, adderall). The bottom of the interface features a text input field with the placeholder "Type your message..." and a "Send" button.

CHARACTERS


Every character is a full configuration

A character in Lagoon is a JSON config — system prompt, character card, model, sampler settings, voice, and memory preferences, all in one place. Switch characters to switch contexts completely.

Character card + system prompt

The **system prompt** sets the model's role and rules. The **character card** is the in-world identity — background, speech patterns, relationships. Both stay separate so you can update one without touching the other.

You can also load a **system context** from a file — paste in a PDF chapter, a wiki article, a code file — and it injects as context every turn.



If you're coming from SillyTavern: character card maps directly. System prompt is your persona instruction. Author's Note is in the Memory section.

Per-character settings: Model & provider · Temperature / Top-P / Rep. penalty / Max tokens · Uncensored mode · Strip thinking · Venice system prompt on/off · Web search · Style Overseer on/off · TTS voice · E2EE · Avatar

MEMORY

Four memory layers. One coherent story.

Most AI chat apps give you a context window and a hard stop. Lagoon stacks four systems so long conversations stay grounded, characters stay consistent, and nothing important gets lost.

Author's Note

Your real-time steering wheel. Write style instructions, POV reminders, scene direction — anything you want the model to act on right now. Reinjects as a system message at a configurable depth from the bottom of the conversation every single turn.

Default depth: 4 messages from end · configurable per session

Lorebook — Anchors

World info that injects only when it's relevant. Create entries with keywords — when those words appear in the last 15 messages, the entry injects automatically. Priority-sorted. Token-budgeted.

The **reveal system** is unique: mark an entry `character_aware: false` and the character won't know that fact yet. The moment they learn it in-scene, Lagoon flips the flag automatically.

Scans last 15 messages · 4,000 token budget · priority-sorted

Semantic Memory — RAG

Lagoon embeds your entire conversation history locally using `all-MiniLM-L6-v2` (no API, no external service). Each turn, it finds the 3 most semantically similar past exchanges and injects them as context.


Context Summaries

When the conversation reaches 75% of the context window, Lagoon summarises the oldest messages in a background thread. You review and approve before anything is pruned. The summary stack grows with the conversation

Fully local · cosine similarity
 ≥ 0.35 · 800 token budget

Always keeps the last 5 full exchange pairs in raw form for freshness.

Anchors modal: Add / edit / delete lore entries · Set keywords, content, priority · Toggle character awareness · Enable/disable per entry



CHAT AREA


Markdown in the chat area

Assistant responses render with full markdown — headers, bold, italics, lists, code blocks, tables. If you're writing a character who produces structured output, or a story with formatted sections, it renders cleanly in the conversation.

What renders

Headings, **bold**, *italics*, - lists, `code blocks`, tables, horizontal rules. Thinking tags (`<think>`) are stripped before display when Strip Thinking is enabled — the reader sees only the response.


The raw text is always preserved in the message history exactly as the model produced it — the markdown rendering is display-only.



WRITING TOOLS

Per-message tools for writers

Every assistant message has an action bar. These aren't generic chat features — they're built around how collaborative fiction actually works.



Fork

Branch the story at any point. Fork creates a new chat with all messages up to and including the selected response. The original is untouched. Explore alternate paths without losing your main thread.

Nudge & Regenerate

Type a behind-the-scenes instruction — `(( more tension, she's scared ))` — and regenerate. The instruction is injected into context for this generation only. It doesn't appear in the chat history. The story doesn't know you interfered.

Edit in Place

Click any assistant message to edit it directly. After you save, Lagoon analyses what you changed — word substitutions, deleted phrases, structural shifts — and stores a style note tagged (noted.) on the message. The model learns from your corrections.



Writing Tools Menu

A context menu of custom rewrite prompts per message. Ships with: Less Tell/More Show, More Sensory Detail, Less Sensory Detail, Add Inner Monologue. Fully customisable — edit labels and prompts, add as many as you want. Stored in your browser.

Regenerate

Retry any response from scratch. Discards the current assistant message and re-runs the model from the same user input. Combine with the Writing Tools menu for targeted regeneration.

Keep for Export

Checkbox on every assistant message. Mark the responses you want to keep. Export compiles only the marked messages. The export count shows in the toolbar. Kept selections survive chat reloads and are saved with the chat file.

STYLE OVERSEER

LLM-powered prose enforcement


The Overseer runs a second LLM after every response and flags violations of your style rules. Accept a correction and it patches the text inline. Dismiss it and move on. It's compounding — accepted corrections inject as notes into the next turn so the model adapts.

Your rules, your model

Six built-in rules ship with Lagoon (no isolated dramatic sentences, strict 3rd person limited, no correction acknowledgment, no unattributed dialogue, etc.). Add your own in the Overseer tab — one rule per line.

The Overseer model is configurable separately from your writing model. Use a fast, cheap model for rules-checking. Use a powerful model for nuanced prose coaching. Preset rule sets for GLM, DeepSeek, Llama, and Qwen are included.

Auto-accept mode: corrections apply automatically without review. Good for absolute bans — banned words never make it to the page.



STYLE OVERSEER IN ACTION — METAPHOR BAN RULE

BEFORE

The offer landed with the precision of a line of perfect code. The silence stretched between them.

✓ AFTER

Kelly's fingers stopped mid-type.
She read the message again.
Then a third time.

🚩 2 violations — hover a message to see this

RULE 4 — ABSOLUTE BAN: METAPHORS

"landed" → "settled"

Replace the metaphor with a concrete verb.

✓ Accept

Dismiss

RULE 5 — ABSOLUTE BAN: SIMILES

"like the wind" → "with long strides"

Replace the simile with a physical description of motion.

✓ Accept

Dismiss

DUAL MODEL

Two characters. Automated conversation.


Pick two characters, write an opening line, hit start. Lagoon runs them in alternating turns — each model fully configured with its own system prompt, sampler settings, and context. They respond to each other. You watch it unfold, or step in at any point.

How it works

Each character uses their full config — their own model, temperature, system prompt, author's note, strip-thinking setting. Character A responds, then character B, back and forth for up to a configurable number of turns.

Controls: Pause mid-conversation, resume, continue past the turn limit, or stop and edit manually. Regenerate or delete any individual exchange. The turn counter tracks where you are.

Good for: generating dialogue drafts, stress-testing character voices against each other, automated back-and-forth RP scenes you can edit into shape afterward.



TRIO MODE

Three-way conversations with emotional intelligence.

Trio Mode extends Dual Model into a full three-way conversation: you and two AI characters, all in the same chat. Characters can interrupt each other, react passively, and track shifting emotional states across turns — producing dialogue that feels alive rather than turn-based.

Session Builder

The Trio Session Builder lets you configure every aspect of a three-way conversation before it starts. Pick two existing characters, define their relationship dynamics, set initial emotional states, and toggle which interactive behaviours are active.

Every field has a hover tooltip explaining what it does — from patience thresholds to ambient mode — so you never have to guess what an option controls.

Trio Session Builder

Relationship Dynamics

Define how each character feels about the other — $A \rightarrow B$ and $B \rightarrow A$ dynamics are injected into each character's system prompt separately. A shared history field gives both characters the same background context, so they reference the same past events without contradiction.

Interactive Behaviours

Interruptions — The non-speaking character can cut in mid-response once the speaker passes their patience threshold (configurable in characters of

dialogue). Lower thresholds produce more impulsive, reactive characters.

Passive Reactions — The non-speaking character produces body language, facial expressions, and muttered asides without taking a full speaking turn. Adds texture and presence to the scene.

Emotional Tracking — Each character's emotional state is tracked across turns and updated based on what happens in the conversation. A character who starts guarded may shift to vulnerable after a confession, and their subsequent responses reflect that shift in tone and word choice.

Ambient Mode — Both characters respond to every user message simultaneously rather than taking turns. Useful for compare-and-contrast scenes where you want both perspectives at once.

Turn Management

Max Consecutive Turns prevents one character from monopolising the conversation — after the limit, the other character is forced to respond.

Let Them Talk mode runs rounds of back-and-forth dialogue between the two characters without user input, then yields control back to you. Set the number of rounds (1–6) in the builder.

Participant Strip

A coloured participant strip at the top of the chat shows both characters, their current emotional state, and who is actively speaking. Trio configs are badged with a **TRIO** tag in the sidebar, and the right-click context menu offers "Edit Trio Session" to reopen the builder and adjust any setting mid-conversation.

EXPORT

Get your writing out cleanly

Mark responses with the Keep checkbox. When you're ready, export compiles only the marked messages — stripped of OOC instructions, system messages, and metadata. Three output formats.

Markdown — clean, portable, `---` separators. Paste into Obsidian, a blog, anywhere.

HTML — self-contained with inline CSS. User and assistant messages styled separately.

Plain text — `[ROLE]: content` format with line separators.

`(( OOC instructions ))` are stripped from user messages automatically.

What exports is the story, not the scaffolding.

PROVIDERS & MODELS

Any model, any endpoint

Lagoon routes to any OpenAI-compatible API. Add models through the interface — provider is inferred automatically from the installed models list.

Venice.ai

Full Venice feature support: uncensored models, web search, TEE models for E2EE, model discovery modal, balance display. Primary supported provider.

Ollama **NO KEY**

Local models at `localhost:11434`. No API key required. Add any model you've pulled with Ollama — it auto-discovers via the install modal.

Custom Endpoint

Any OpenAI-compatible API: KoboldCpp, LM Studio, llama.cpp, text-generation-webui, or any hosted endpoint. Set a base URL and optional API key per endpoint.

Google Gemini

Gemini Live for real-time voice sessions. Text models via the standard endpoint. Requires a Google API key.

VOICE

TTS and Gemini Live

Read responses aloud with per-character voice settings, or go fully hands-free with Gemini Live real-time voice sessions.

Text-to-Speech

Two TTS providers: **Venice** (23 voices including af_sky, af_alloy, am_adam) and **Google** (30 voices including Aoede, Charon, Puck, Fenrir). Set a default voice per character — it triggers automatically when auto-read is on, or manually via the speaker button on any message.

Gemini Live

Real-time bidirectional voice via Google's Gemini Live API. Speak to your character, hear them respond in near real-time. Custom system prompt per session. 30 voice options. Enabled from the mobile drawer or desktop voice panel.

PRIVACY

End-to-end encryption for sensitive RP

Venice TEE (Trusted Execution Environment) models run in hardware-isolated enclaves. Lagoon's E2EE support encrypts your messages before they leave your machine.

When E2EE is enabled on a character config using a Venice TEE model, Lagoon performs a **secp256k1 ECDH key exchange** with the TEE at session start, derives a shared key via **HKDF-SHA256**, and encrypts all messages with **AES-256-GCM** before transmission. The server never sees plaintext. Attestation is verified (signed, non-debug, nonce-matched) before any keys are exchanged. Venice's own infrastructure cannot read the conversation.

Lagoon — Local-first AI writing workspace

[GitHub](#)